



# Airbnb Price Prediction

Springboard Bootcamp 2021

Alena Tskhondiya

# Project Background and Aim

Airbnb is an internet marketplace for short-term home and apartment rentals. It allows you to rent out your home for a week while you're away or rent out your empty bedroom. One challenge that Airbnb hosts face is determining the optimal nightly rent price.

For this project, their data set scraped on April 1, 2021, on the city of Seattle, WA, USA is used. It contains information on all Seattle Airbnb listings that were live on the site on that date (over 4,200). Here's a [direct link](#).

Each row in the data set is a listing available for rental in Airbnb's site for the specific city (observations). The columns describe different characteristics of each listing (features).

- accommodates
- bedrooms
- Bathrooms
- Beds
- Price
- minimum\_nights
- maximum\_nights
- number\_of\_reviews

To model the spatial relationship between Airbnb rental prices and property proximity to certain venues, we use the [Foursquare API](#) to access the city's venues and the street network, available through [OpenStreepMap \(OSM\)](#).

# Table of Contents

**01**    **Cleaning and Pre-processing**

**02**    **Exploratory Data Analysis**

**03**    **Walkability to nearest venues**

**04**    **Preparing data for modelling**

**05**    **Modelling**

01

# **Cleaning and Pre-processing**

# Initial data. Dropping columns

4213 rows, 74 columns

40 columns – dropped

## Dropped:

- text columns have been dropped since Natural Language Processing will not be used in the creation of this model.
- Columns with several NULL entries are dropped
- Columns with the same null cases are dropped and will keep one column
- multiple columns for property location can be dropped and one column for area will be kept, neighborhood\_cleansed.
- Two main columns will be used minimum\_nights and maximum\_nights
- Checking the boolean and categorical features, and the one that contains one category can be dropped

# Cleaning individual columns

- host\_since
- host\_response\_time
- host\_response\_rate
- host\_is\_superhost
- host\_listings\_count
- host\_identity\_verified
- neighbourhood\_cleansed
- property\_type
- room\_type
- accommodates
- bathrooms
- bedrooms
- beds
- amenities
- Price
- minimum\_nights
- maximum\_nights
- availability\_30
- availability\_60
- availability\_90
- availability\_365
- number\_of\_reviews
- number\_of\_reviews\_ltm
- first\_review
- last\_review
- review\_scores\_rating
- review\_scores\_accuracy
- review\_scores\_cleanliness
- review\_scores\_checkin
- review\_scores\_communication
- review\_scores\_location
- review\_scores\_value
- instant\_bookable
- reviews\_per\_month

# Cleaning individual columns

- **host\_since.** This datetime column will be converted into a measure of the number of days that a host has been on the platform, measured from the date that the data was scraped
- **host\_response\_time** has 18% unknown listings, it will be retained as its own category, 'unknown'
- **host\_response\_rate** 70% of hosts respond 100% of the time, this will be kept as its own category, and other values will be grouped into bins
- **host\_is\_superhost** There are 192 row with no values for each of three different host-related features. These rows will be dropped
- **property\_type** The categories `Apartment`, `House` and `Other` will be used, as most properties can be classified as either apartment or house.
- **bathrooms, bedrooms and beds** Missing values will be replaced with the median (to avoid strange fractions)
- **amenities** will be extracted based on quick research into which amenities are considered by guests a selection of the more important as well as personal experience. Amenity features contains fewer than 10% of listings will be removed
- **availability** There are multiple different measures of availability, which will be highly correlated with each other. Only one will be retained, availability for 365 days
- **first\_review and last\_review** Almost 20 percent of listings have not had a review written for them. This is too large a proportion of the dataset to drop and replace with median/mean values. These will be kept as an `unknown` category, and the feature will have to be treated as categorical (and therefore one-hot encoded) rather than numerical.
- **review ratings columns** The listings without reviews will be kept and replaced with `unknown`. Other ratings will be grouped into bins.
- **number\_of\_reviews\_ltm and reviews\_per\_month** These will be highly correlated with `number\_of\_reviews` and `reviews\_per\_month` and so will be dropped.

02


# **Exploratory Data Analysis**



# Time Series


Time is an important factor to consider in a model when we wish to predict prices or trends. There are questions that come to play when dealing with time series. For example: Is there any seasonality to the price? Is it stationary? Even though we are not going to include this aspect into our model, it is good to explore it to be aware of it and be able to make recommendations for future research.

Seattle hosts joining Airbnb and listings getting their first review in each month




# Time Series

Every year, you see a peak towards hosts joining around the middle of the year (summer), and the lowest points are the beginning and the end of each year. There is a big peak in the number of hosts joining Airbnb between 2013 and 2014. Indeed, there has been a fast growth of Airbnb since middle 2013.




# Time Series

Number of Airbnb listings getting their first review each month




# Time Series

There are a number of professional Airbnb management companies which host a large number of listings under a single host profile. However, there is no consistent upwards trend in the average number of properties managed by each host.



# Time Series


In term of changes in prices over time, the average price per night for Airbnb listings in Seattle has increased slightly over the last 10 years. In particular, the top end of property prices has increased, resulting in a larger increase in the mean price compared to the median. The mean price in 2010 was 119.29 and the median 73.0, whereas the mean price in 2020 (the last complete year of data) was 122 and the median 104.



# Numerical Features


The most common property setup sleeps two people in one bed in one bedroom, with one bathroom. Unsurprisingly, properties that accommodate more people achieve noticeably higher rates per night, with diminishing returns coming after about 10 people.

Median price of Airbnbs accommodating different number of guests




# Categorical Features

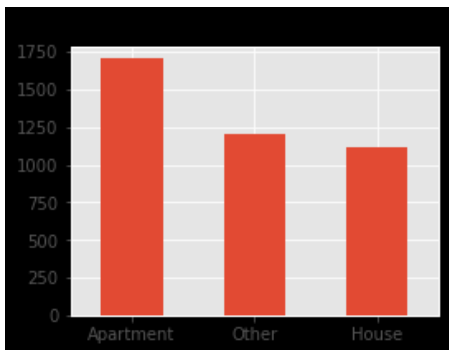
Number of Airbnb listings in each Seattle borough



Median price of Airbnb listings in each Seattle borough

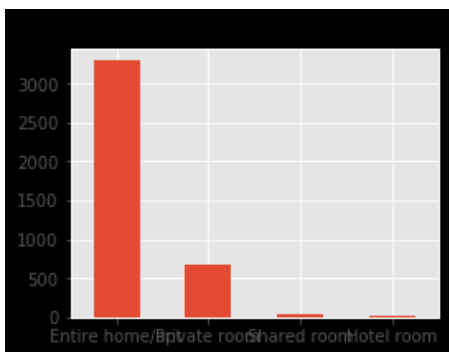


# Categorical Features




## Property type

Apartment 0.423278  
Other 0.299925  
House 0.276797



## Room type

Entire home/apt 0.818702  
Private room 0.166377  
Shared room 0.011440  
Hotel room 0.003482



## Overall listing rating

No review - 798  
95-100% - 2317  
80-94% - 882  
0-79% - 125




03

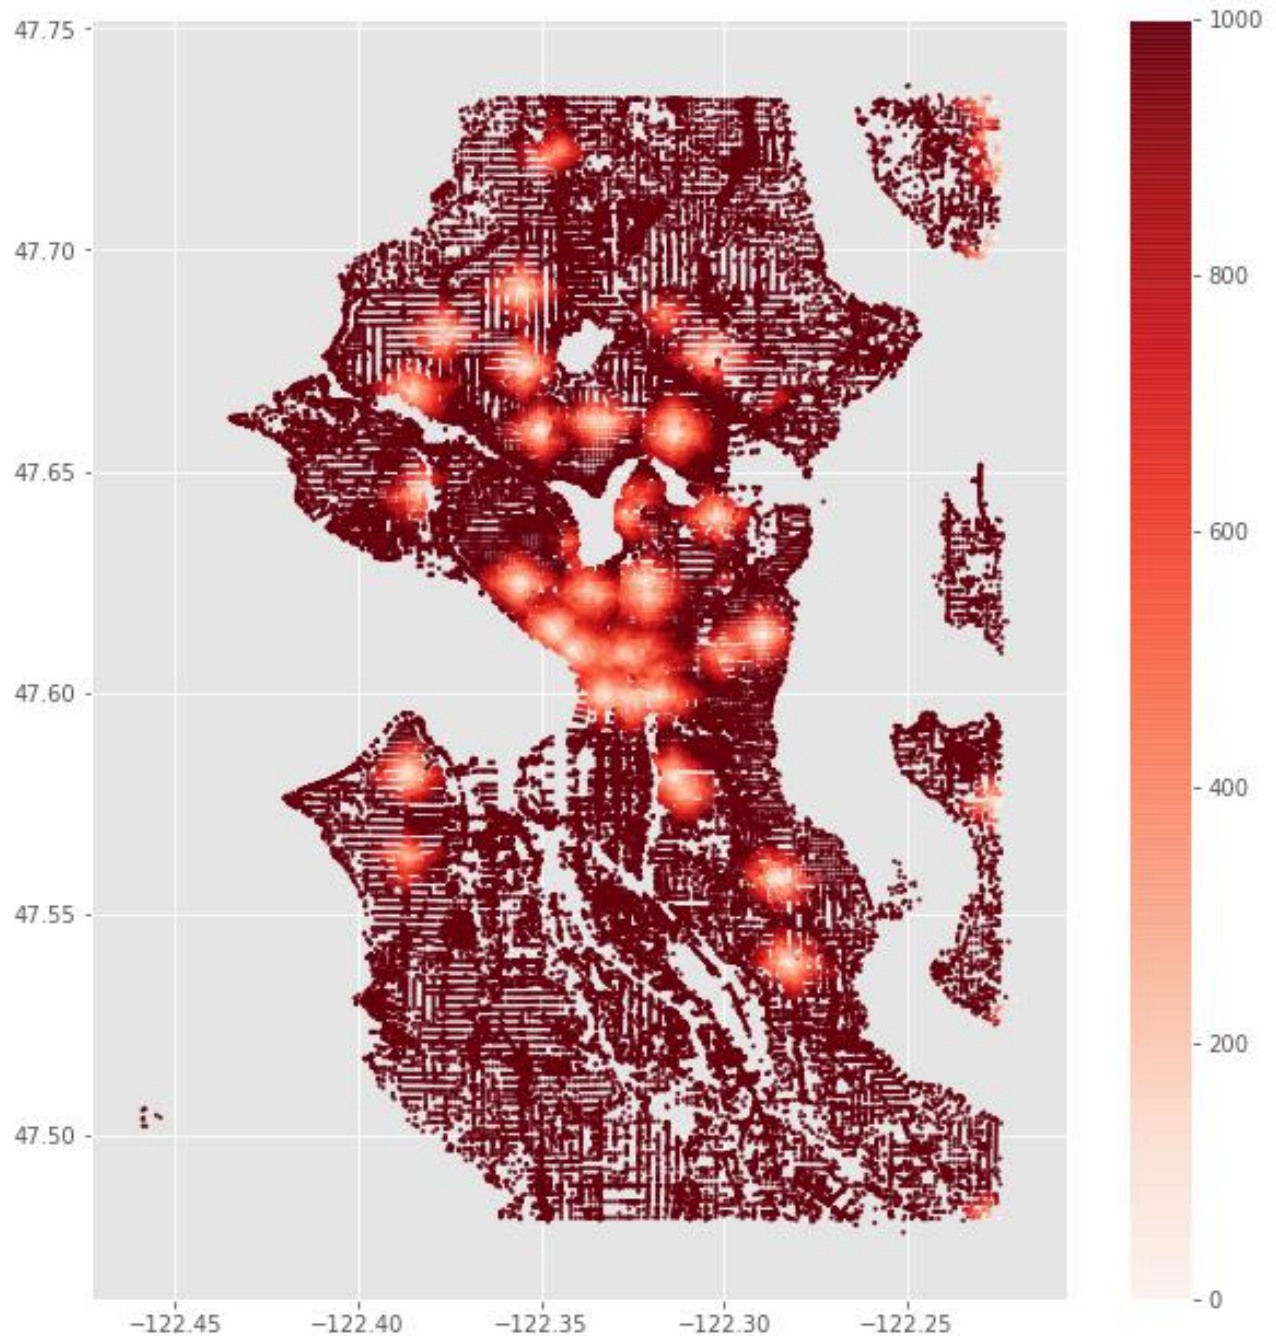
# **Walkability to nearest venues**

# Walking distance (m) to nearest amenity around Seattle

As part of our model, we are looking to explore proximity to certain venues as a possible price predictor. Walkability and ability to reach places (touristic attractions, restaurants, cafes and shops) maybe a deal-maker or breaker when it comes to choosing an accommodation. For this, we will use Foursquare API to explore the venues per neighborhood.

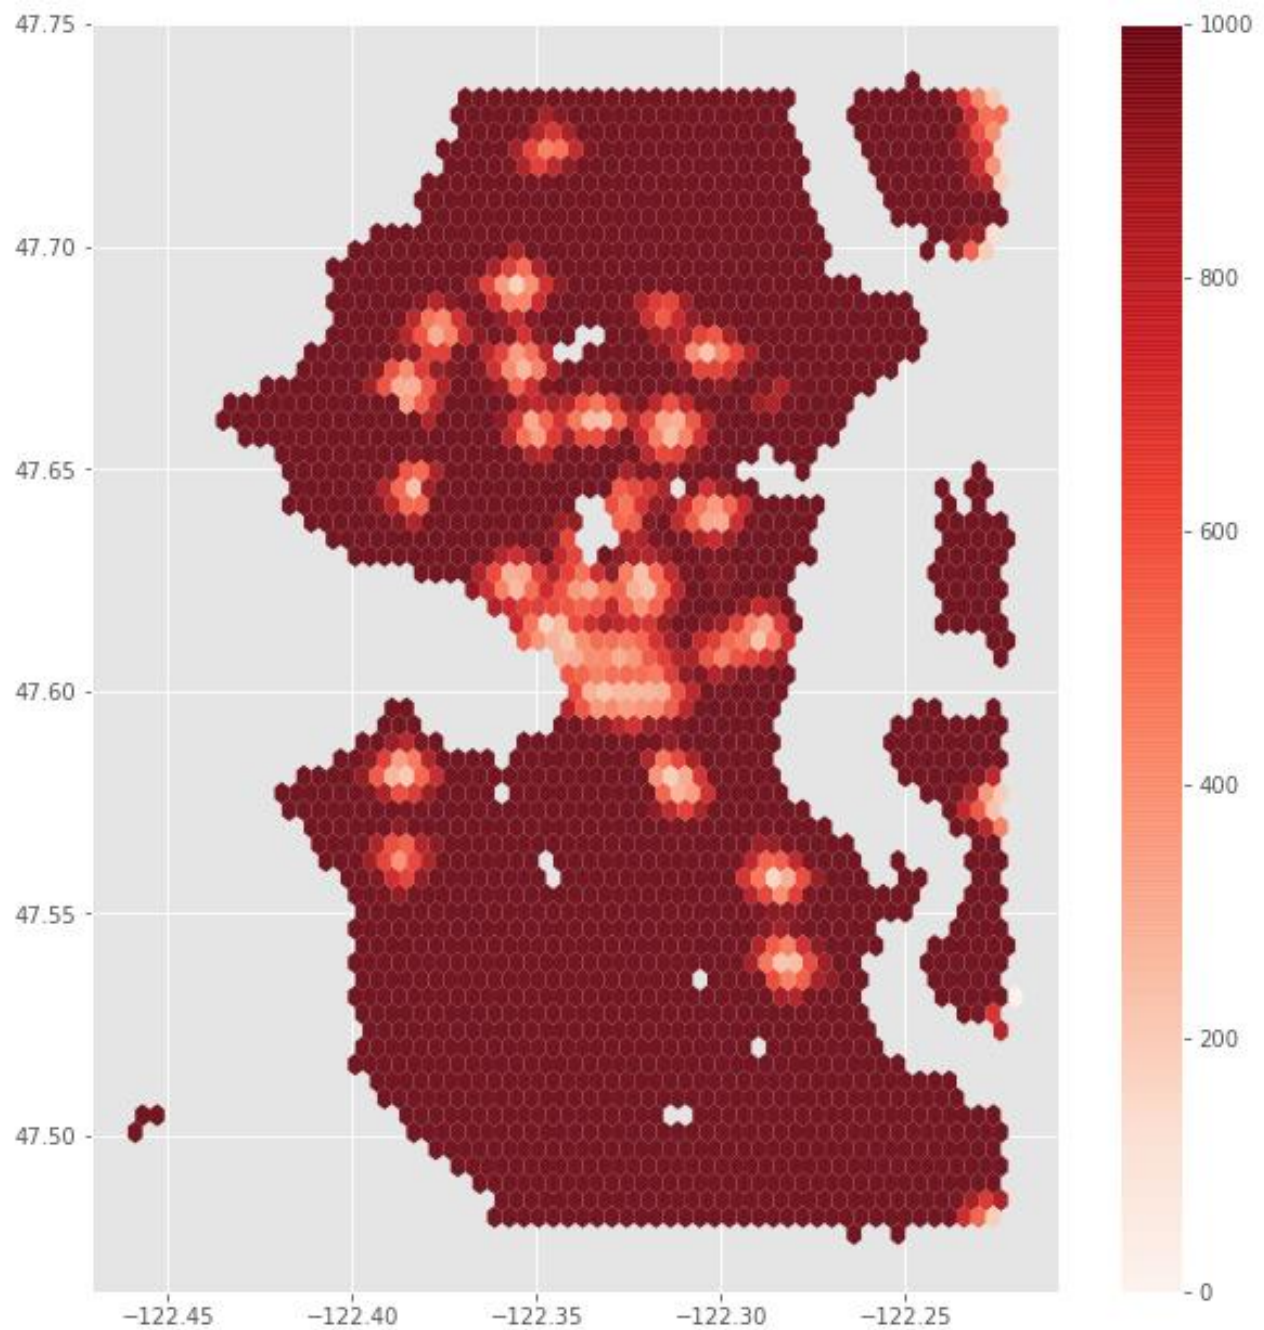


# Walking distance (m) to 5th nearest amenity around Seattle






# Walking distance (m) to nearest amenity around Seattle




# Walking distance (m) to 5th nearest amenity around Seattle

Access to restaurants, shops, cafes, bars and pubs are taken as essential for Airbnb users. All amenities are waited equally and using distance to the fifth nearest amenity as a compound measure of accessibility. This gives a clearer picture of which neighborhoods are most walkable, compared with plotting just the distance to the single nearest venue/amenity.



# Network aggregation. Score

The compound measure of accessibility (network distance from the node to the 5th nearest POIs).



04

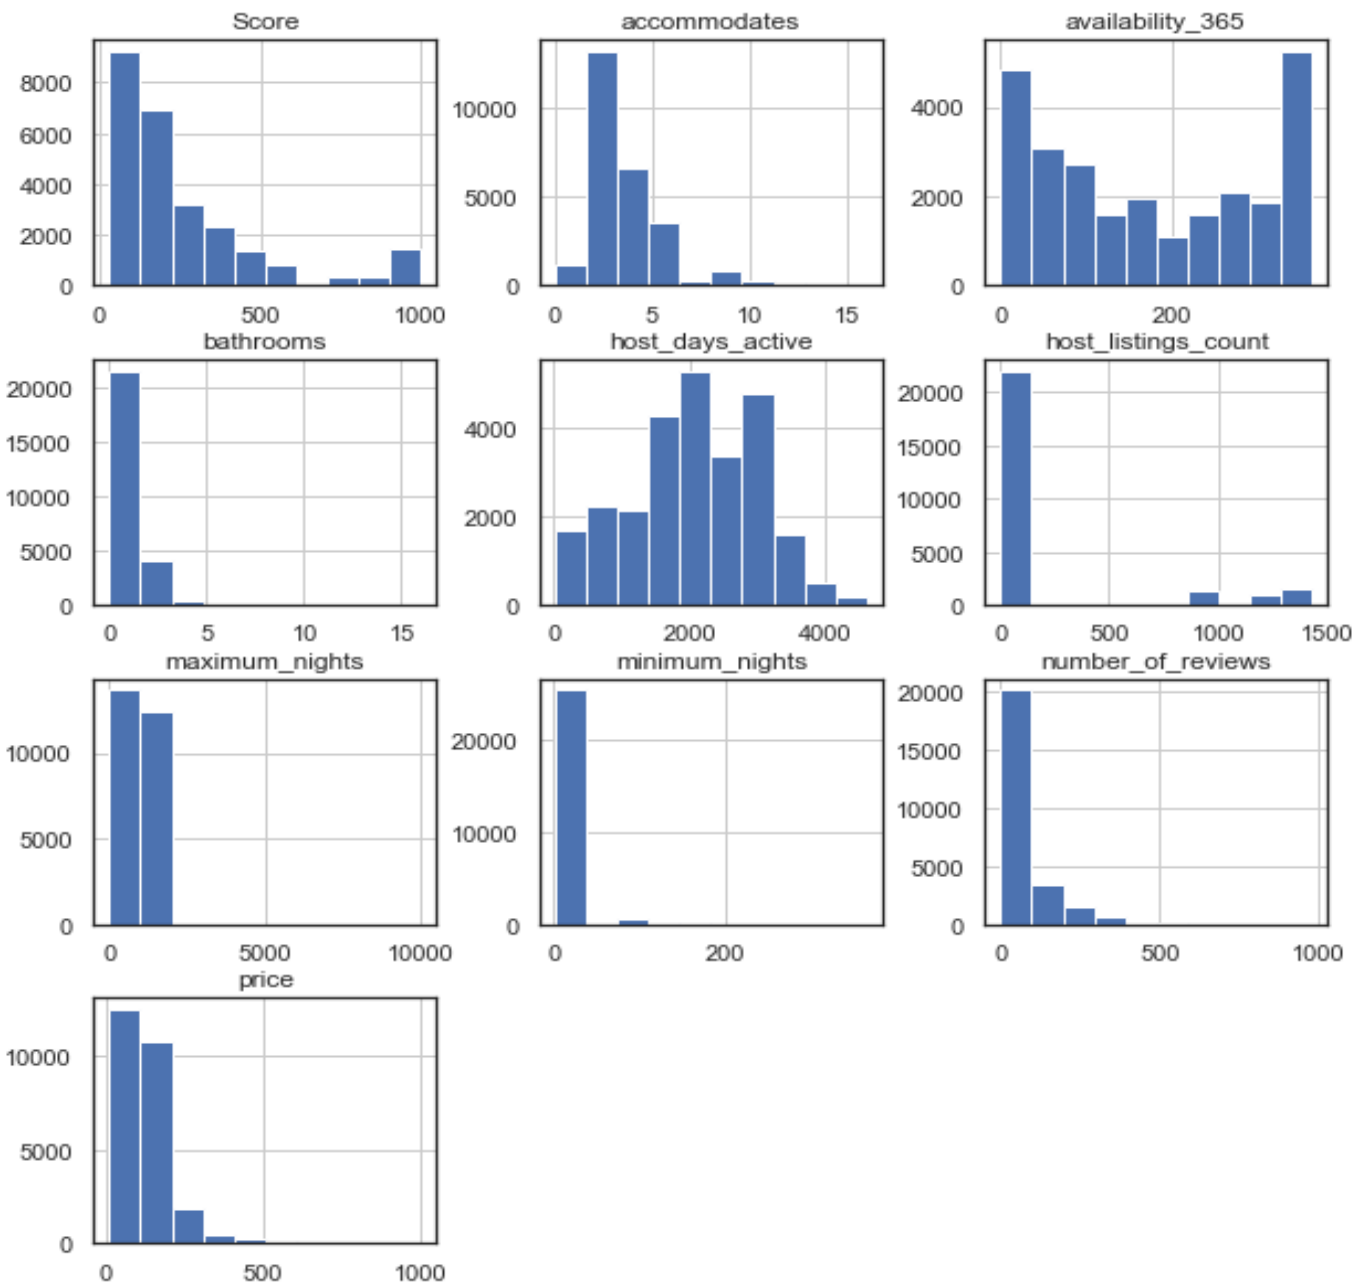
# Preparing data for modelling

# Multicollinearity Heatmap

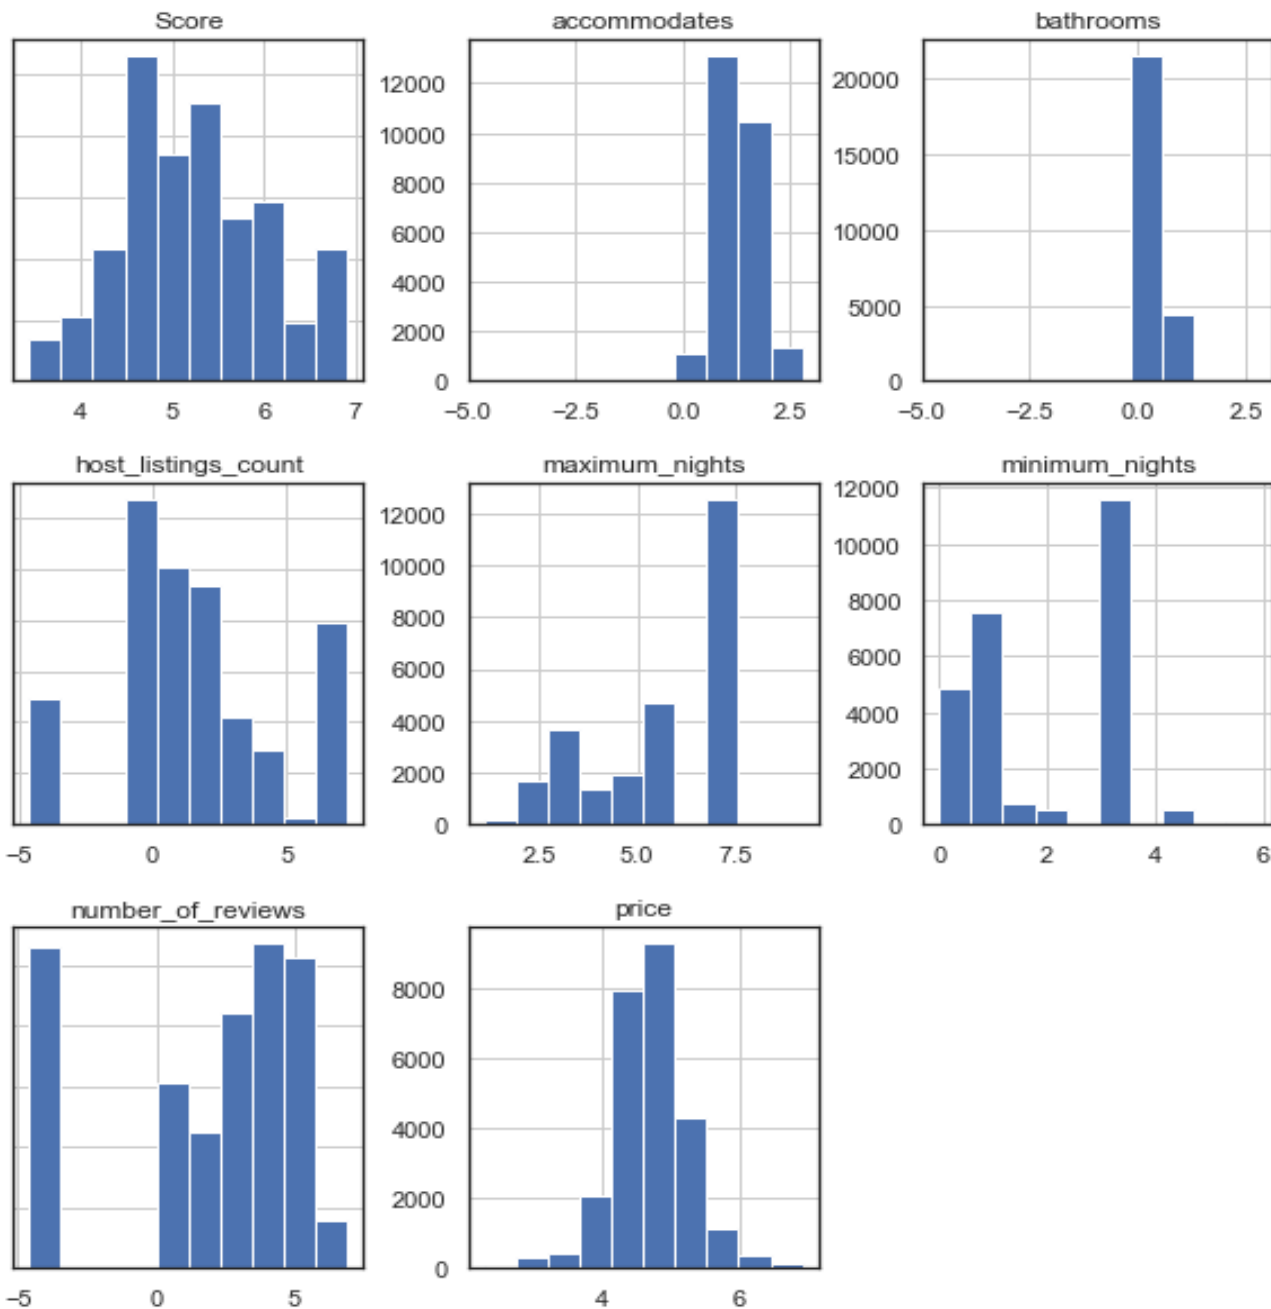




# Standardizing and normalizing



# Standardizing and normalizing




The background features a bold, geometric design. The top half is primarily yellow, with a grey triangular shape on the left side. The bottom half is white, with a grey horizontal bar on the left and a yellow triangular shape on the right. The number '05' is prominently displayed in white on the yellow background.

05

# Modelling

# Feature importance

### Feature importances in the XGBoost model



# Feature importance

Weight

room_type_Entire home/apt	0.418624
bathrooms	0.063108
gym	0.028660
elevator	0.021742
review_scores_rating_no reviews	0.018268
tv	0.014842
property_type_Other	0.014568
accommodates	0.013613
Neighbourhood_Pike-Market	0.012118
Neighbourhood_South Lake Union	0.011350

# Models

01

**Spatial Hedonic Price Model (HPM)**

Training RMSE: 0.1018  
Validation RMSE: 0.0991  
Training r2: 0.6729  
Validation r2: 0.6672

02

**Gradient boosted decision trees**

Training MSE: 0.0054  
Validation MSE: 0.0093  
Training r2: 0.9825  
Validation r2: 0.9688

03

**Hedonic regression with dropped columns**

Training RMSE: 0.1018  
Validation RMSE: 0.0991  
Training r2: 0.6729  
Validation r2: 0.6672

04

**XG Boost with dropped columns**

Training MSE: 0.0054  
Validation MSE: 0.0093  
Training r2: 0.9825  
Validation r2: 0.9688

# Conclusions

The best performing model is the XGBoost model (Model 4) was able to predict 98.25% of the variation in price with an RMSE of 0.01. Which means we still have a remaining 34.44% unexplained. This could be due to several other features that are not part of our dataset or the need to analyze our features more closely.

For example, given the importance of customer reviews of the listing in determining price, perhaps a better understanding of the reviews could improve the prediction. Using Sentiment Analysis, a score between -1 (very negative sentiment) and 1 (very positive sentiment) can be assigned to each review per listing property. The scores are then averaged across all the reviews associated with that listing and the final scores can be included as a new feature in the model.

It was noticeable that reviews about listing location, rather than the location features themselves, were higher in the feature importance list. Thus, this finding could perhaps be used by Airbnb hosts when writing their listing's description. Highlighting accessibility and location benefits of staying with them could perhaps benefit them and how much they can ask for their listing.

# Thanks

Do you have any questions?  
youremail@freepik.com  
+91 620 421 838  
yourcompany.com



CREDITS: This presentation template was  
created by **Slidesgo**, including icons by **Flaticon**,  
and infographics & images by **Freepik**

**Please keep this slide for attribution**



# Instructions for use

In order to use this template, you must credit [Slidesgo](#) by adding the credits line “This template was created by Slidesgo”.

**You are allowed to:**

- Modify this template.
- Use it for both personal and commercial projects.

**You are not allowed to:**

- Sublicense, sell or rent any of Slidesgo Content (or a modified version of Slidesgo Content).
- Distribute Slidesgo Content unless it has been expressly authorized by Slidesgo.
- Include Slidesgo Content in an online or offline database or file.
- Offer Slidesgo templates (or modified versions of Slidesgo templates) for download.
- Acquire the copyright of Slidesgo Content.

For more information about editing slides, please read our FAQs or visit Slidesgo School:

<https://slidesgo.com/faqs> and <https://slidesgo.com/slidesgo-school>

# Fonts & colors used

This presentation has been made using the following fonts:

## **Roboto**

(<https://fonts.google.com/specimen/Roboto?query=ro>)

#fccc3b

#606969

#2e383d

#f3f1ea

# Stories by Freepik

Create your Story with our illustrated concepts. Choose the style you like the most, edit its colors, pick the background and layers you want to show and bring them to life with the animator panel! It will boost your presentation. Check out [How it Works](#).



Pana



Amico



Bro



Rafiki



Cuate

